

Analyzing cross-over trials with the package LMMstar

Brice Ozenne

October 22, 2025

In the context of a cross-over trial, this vignette discusses

- (i) how mixed model generalizes t-tests
- (ii) modeling choices regarding the variance-covariance structure implied by the choice of the repetition variable (period or treatment).

We will use to the following  packages:

```
library(LMMstar)
library(ggplot2)
```

1 Illustrative dataset

The `bloodpressureL` dataset:

```
data(bloodpressureL, package = "LMMstar")
head(bloodpressureL)
```

	<code>id</code>	<code>sequence</code>	<code>treatment</code>	<code>period</code>	<code>duration</code>
1	1	ABC	A	1	1.9
2	1	ABC	B	2	2.9
3	1	ABC	C	3	4.3
4	2	ABC	A	1	1.4
5	2	ABC	B	2	2.3
6	2	ABC	C	3	3.0

originates from a cross-over trial comparing the impact of three formulations of a drug on the blood pressure. The study was conducted on 12 male volunteers randomly divided into three groups (`sequence`) and receiving each of the three formulations (`treatment`) with a wash-out period of one week. The outcome variable is `duration` where a larger duration indicate a better outcome as the blood pressure is under control for a longer time. While there are not missing values in this dataset:

```
sum(is.na(bloodpressureL))
```

```
[1] 0
```

only 3 out of the 6 possible sequences of treatment have been allocated:

```
levels(bloodpressureL$sequence)
```



```
[1] "ABC" "BCA" "CAB"
```

A spaghetti plot provides a graphical representation of the dataset, with either period or treatment on the x-axis:

```
ggTime <- ggplot(bloodpressureL, aes(x = period, y = duration, group = id))  
ggTime <- ggTime + geom_line() + geom_point(aes(color = treatment), size = 2)  
ggTime
```

or

```
ggTreat <- ggplot(bloodpressureL, aes(x = treatment, y = duration, group = id))  
ggTreat <- ggTreat + geom_line() + geom_point(aes(color = period), size = 2)  
ggTreat
```



2 Matching the t-test results

One can use a paired t-test to assess the treatment effect when there is not missing values and no covariate to adjusted on (in particular no period effect). It is easier to carry-out with the wide format of the dataset:

```
bloodpressureW <- reshape(bloodpressureL, direction = "wide",
                           idvar = "id", timevar = "treatment",
                           v.names = c("duration", "period"))
head(bloodpressureW)
```

	id	sequence	duration.A	period.A	duration.B	period.B	duration.C	period.C
1	1	ABC	1.9	1	2.9	2	4.3	3
4	2	ABC	1.4	1	2.3	2	3.0	3
7	3	ABC	1.4	1	2.8	2	3.5	3
10	4	ABC	0.6	1	2.1	2	2.9	3
13	5	BCA	2.2	3	2.2	1	3.6	2
16	6	BCA	1.3	3	2.1	1	2.7	2

For instance we can compare drug B and A using:

```
t.test(bloodpressureW$duration.B - bloodpressureW$duration.A)
```

One Sample t-test

```
data: bloodpressureW$duration.B - bloodpressureW$duration.A
t = 2.9, df = 11, p-value = 0.015
alternative hypothesis: true mean is not equal to 0
95 percent confidence interval:
 0.13805 1.01195
sample estimates:
mean of x
 0.575
```

To retrieve the same results with a linear mixed model, one can use treatment as indexing the repetitions, i.e., model a treatment specific variance and correlation:

```
eTreat.lmm2tt <- lmm(duration ~ treatment, repetition = ~treatment|id, data = bloodpressureL)
model.tables(eTreat.lmm2tt)
```

	estimate	se	df	lower	upper	p.value
(Intercept)	1.7250	0.16703	11.002	1.35739	2.0926	5.3415e-07
treatmentB	0.5750	0.19853	10.998	0.13804	1.0120	1.4542e-02
treatmentC	1.2583	0.22578	10.998	0.76137	1.7553	1.6701e-04

⚠ using period as the repetition variable, i.e., modeling a period specific variance and correlation, would lead to different estimates:

```
ePeriod.lmm2tt <- lmm(duration ~ treatment, repetition = ~period|id, data = bloodpressureL)
model.tables(ePeriod.lmm2tt)
```

	estimate	se	df	lower	upper	p.value
(Intercept)	1.68755	0.20349	4.7145	1.15478	2.2203	5.5048e-04
treatmentB	0.58766	0.19895	14.4584	0.16223	1.0131	1.0173e-02
treatmentC	1.16557	0.19654	11.9624	0.73718	1.5939	7.0104e-05

As shown in appendix A.2, this mixed model considers both the variable treatment and period when deciding how much each observation contributes to the estimation of a given parameter. On one side, it makes sense that an observation taken at a period with large variance should contribute less to parameter estimation compared to an observation taken at a period with low variance. On the other side, it can be suprising that treatment B outcomes can contribute to the estimation of treatment A. This is however not the case in absence of period effects since the weights sum to 0 for treatment B and C when estimating the intercept.

⚠ using a random intercept model instead would lead to the same estimate but a different p-value:

```
eTreat.RI <- lmm(duration ~ treatment + (1|id), data = bloodpressureL)
model.tables(eTreat.RI)
```

	estimate	se	df	lower	upper	p.value
(Intercept)	1.7250	0.15192	29.483	1.41452	2.03548	2.7569e-12
treatmentB	0.5750	0.18673	22.000	0.18774	0.96226	5.4846e-03
treatmentC	1.2583	0.18673	22.000	0.87107	1.64560	8.9931e-07

as it makes more restrictive assumptions (homoschedasticity, equal correlation).

3 Accounting for a period effect

A natural extension of the t-test to adjust for a possible period effect on the average outcome is to consider the corresponding mixed model (i.e. treatment as repetition) and add period in the mean model:

```
eTreat.lmm <- lmm(duration ~ treatment + period, repetition = ~treatment|id,  
                  data = bloodpressureL)  
model.tables(eTreat.lmm)
```

	estimate	se	df	lower	upper	p.value
(Intercept)	1.54915	0.16604	13.9154	1.192831	1.90547	2.3014e-07
treatmentB	0.57500	0.16768	9.3858	0.198043	0.95196	7.0688e-03
treatmentC	1.25833	0.17889	9.2795	0.855515	1.66115	5.2215e-05
period2	0.19991	0.12698	4.0388	-0.151307	0.55112	1.8984e-01
period3	0.32764	0.12109	4.8784	0.014028	0.64125	4.3588e-02

Here because the design is balanced in term of period across treatments, we obtain the same estimates for the difference in treatment effect as if we do not adjust for period. However the estimated mean outcome under each treatment (say treatment A) now depends on all observations (and not only observations under treatment A). See appendix [A.3](#) for details.

4 What if there is a baseline measurement?


Consider now another study design where all patients have a baseline measurement before receiving each treatment. As an illustrative example we will consider the following illustrative dataset:

```
rho <- c(AB = 0.3, bb = 0.9, bA = 0.7, bB = 0.6)
SigmaBC0 <- rbind(cbind(matrix(c(1, rho["bA"], rho["bA"], 1), 2, 2),
                             matrix(c(rho["bb"], rho["bA"], rho["bB"], rho["AB"]), 2, 2)),
                  cbind(matrix(c(rho["bb"], rho["bB"], rho["bA"], rho["AB"]), 2, 2),
                             matrix(c(1, rho["bB"], rho["bB"], 1), 2, 2)))
muBC0 <- c(b1 = 0, A = 1, b2 = 0, B = 1.5)

library(mvtnorm)
set.seed(10)
n.obs <- 15
M1 <- data.frame(id = 1:n.obs, sequence = "AB",
                 rmvnorm(n.obs, mean = muBC0, sigma = SigmaBC0))
names(M1)[3:6] <- paste0("T", 1:4)
M2 <- data.frame(id = n.obs+(1:n.obs), sequence = "BA",
                 rmvnorm(n.obs, mean = muBC0[c(1, 4, 3, 2)],
                        sigma = SigmaBC0[c(1, 4, 3, 2), c(1, 4, 3, 2)]))
names(M2)[3:6] <- paste0("T", 1:4)
dfL.BC0 <- reshape(rbind(M1, M2), direction = "long",
                  idvar = "id", varying = names(M1)[-(1:2)], v.names = c("Y"), times = 1:4)
dfL.BC0$treatment <- "baseline"
dfL.BC0$treatment[dfL.BC0$time == 2 & dfL.BC0$sequence == "AB"] <- "A"
dfL.BC0$treatment[dfL.BC0$time == 4 & dfL.BC0$sequence == "AB"] <- "B"
dfL.BC0$treatment[dfL.BC0$time == 2 & dfL.BC0$sequence == "BA"] <- "B"
dfL.BC0$treatment[dfL.BC0$time == 4 & dfL.BC0$sequence == "BA"] <- "A"
dfL.BC0$treatment <- factor(dfL.BC0$treatment, levels = c("baseline", "A", "B"))
dfL.BC0$period <- as.character(1 + (dfL.BC0$time %in% 3:4))
dfL.BC0[dfL.BC0$id==1,]
```

	id	sequence	time	Y	treatment	period
1.1	1	AB	1	-0.84169	baseline	1
1.2	1	AB	2	0.36197	A	1
1.3	1	AB	3	-1.28127	baseline	2
1.4	1	AB	4	0.57369	B	2

```
gg.BC0 <- ggplot(dfL.BC0, aes(x=time, y = Y, group = id))
gg.BC0 <- gg.BC0 + geom_line()
gg.BC0 <- gg.BC0 + geom_point(aes(color = treatment, shape = treatment), size = 3)
gg.BC0 <- gg.BC0 + facet_wrap(~sequence, labeller = label_both)
gg.BC0
```



4.1 First time period

If we restrict the dataset to the first period (time 1 and 2) we obtain a standard 2 arm randomized trial. To compare mixed model results with t-test tests and ANCOVA, we will use the wide format of the dataset:

```
dfW.BC0 <- reshape(dfL.BC0[c("id", "time", "treatment", "Y")], direction = "wide",
  timevar = "time", idvar = "id", v.names = c("Y", "treatment"))
```

4.1.1 Change

Performing a two-sample t-test on the change:

```
tt.change <- with(dfW.BC0, t.test(y = Y.2[treatment.2=="B"]-Y.1[treatment.2=="B"],
  x = Y.2[treatment.2=="A"]-Y.1[treatment.2=="A"]))
c(estimate = unname(diff(tt.change$estimate)),
  se = tt.change$stderr,
  lower = tt.change$conf.int[1],
  upper = tt.change$conf.int[1],
  df = tt.change$parameter,
  p.value = tt.change$p.value)
```

estimate	se	lower	upper	df.df	p.value
0.529437	0.224916	-0.990797	-0.990797	27.165129	0.026059

is equivalent to a linear mixed model iwht interaction between:

```
e.lmm <- lmm(Y ~ time*sequence, repetition =~time|id,
  data = dfL.BC0[dfL.BC0$time %in% 1:2,])
model.tables(e.lmm)
```

	estimate	se	df	lower	upper	p.value
(Intercept)	-1.59683	0.35438	28.004	-2.322738	-0.87092	1.0687e-04
time	1.16038	0.15904	27.995	0.834599	1.48616	6.0708e-08
sequenceBA	-0.12004	0.50117	28.004	-1.146632	0.90655	8.1244e-01
time:sequenceBA	0.52944	0.22492	27.995	0.068714	0.99016	2.5827e-02

4.1.2 Ancova

Controlling for baseline difference using an ANCOVA:

```
e.ANCOVA <- lm(Y.2 ~ Y.1 + treatment.2, data = dfW.BCO)
summary(e.ANCOVA)$coef
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	1.07449	0.15983	6.7226	3.2294e-07
Y.1	0.80321	0.10762	7.4631	4.9918e-08
treatment.2B	0.61000	0.22050	2.7664	1.0103e-02

has the same estimate of a linear mixed model with baseline constraint:

```
e0.lmm <- lmm(Y ~ treatment, repetition =~time|id,
              data = dfL.BCO[dfL.BCO$time %in% 1:2,])
summary(anova(e0.lmm, effects = c("treatmentB-treatmentA=0")))
```

Hypothesis-specific Wald test

	estimate	se	df	lower	upper	p.value
treatmentB-treatmentA=0	0.61	0.217	27.8	0.166	1.054	0.00882 **

: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1.

df: Satterthwaite approximation w.r.t. model-based se.

se: based on the observed information (model-based).

⚠ When fitting the mixed model, the variance was on purpose modeled to be time dependent instead of treatment dependent to match the ANCOVA. In many applications, however, a treatment dependent variance and correlation is preferable:

```
e0T.lmm <- lmm(Y ~ treatment, repetition =~treatment|id,
               data = dfL.BCO[dfL.BCO$time %in% 1:2,])
summary(anova(e0T.lmm, effects = c("treatmentB-treatmentA=0")))
```

Hypothesis-specific Wald test

	estimate	se	df	lower	upper	p.value
treatmentB-treatmentA=0	0.602	0.216	27.7	0.159	1.046	0.00954 **

: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1.

df: Satterthwaite approximation w.r.t. model-based se.

se: based on the observed information (model-based).

Covariates: to retrieve the same point estimate between the ANCOVA and the mixed model, covariates should be included in the mixed model with a time interaction.

4.2 Multiple time periods

4.2.1 Change

A straightforward analysis with multiple time periods is to evaluate the within-individual difference between treatment of the change from baseline:

```
dfW.BC0$dY <- (dfW.BC0$Y.4-dfW.BC0$Y.3)-(dfW.BC0$Y.2-dfW.BC0$Y.1)
tt.change2 <- with(dfW.BC0, t.test(c(dY[treatment.4 == "B"],-dY[treatment.4 == "A"])))
c(estimate = unname(tt.change2$estimate),
  se = tt.change2$stderr,
  lower = tt.change2$conf.int[1],
  upper = tt.change2$conf.int[1],
  df = tt.change2$parameter,
  p.value = tt.change2$p.value)
```

estimate	se	lower	upper	df.df	p.value
0.479266	0.181401	0.108260	0.108260	29.000000	0.013144

One can retrieve this result by introducing 2 design variables:

```
dfL.BC0$treated <- (dfL.BC0$treatment!="baseline")
dfL.BC0$periodB <- with(dfL.BC0, (period==1 & sequence=="BA") | (period==2 & sequence=="AB"))
e.lmm2tt <- lmm(Y ~ sequence:treated+(treated*periodB),
  repetition = ~time|id, data = dfL.BC0)
model.tables(e.lmm2tt)
```

	estimate	se	df	lower	upper	p.value
(Intercept)	-0.43409	0.264062	28.838	-0.97429	0.106105	1.1105e-01
treatedTRUE	1.14741	0.140692	35.009	0.86179	1.433027	1.3152e-09
periodBTRUE	-0.06553	0.071101	28.992	-0.21095	0.079889	3.6432e-01
sequenceBA:treatedFALSE	0.47492	0.370826	28.037	-0.28463	1.234483	2.1079e-01
sequenceBA:treatedTRUE	0.52510	0.325697	28.013	-0.14205	1.192241	1.1812e-01
treatedTRUE:periodBTRUE	0.47927	0.181401	28.988	0.10825	0.850279	1.3146e-02

where the coefficient of interest in the last line, among the treatment periods (A or B) whether there is a difference between being under treatment B or treatment A.

4.2.2 Ancova

A more 'natural' linear mixed model would simply adjust the treatment effect on period and sequence effect:

```
eBCO.lmm <- lmm(Y ~ period + sequence + treatment, repetition =~time|id,
               data = dfL.BCO)
summary(anova(eBCO.lmm, effects = c("treatmentB-treatmentA=0")))
```

Hypothesis-specific Wald test

```
              estimate      se   df lower upper p.value
treatmentB-treatmentA=0    0.436 0.168 27.4 0.091 0.781  0.0153 *
-----
:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1.
df: Satterthwaite approximation w.r.t. model-based se.
se: based on the observed information (model-based).
```

but this does not match the average individual difference adjusted for both baselines:

```
dfW.BCO$dY <- with(dfW.BCO, (Y.4-Y.2)*(treatment.4=="B")+(Y.2-Y.4)*(treatment.2=="B"))
summary(lm(dY ~ Y.1 + Y.3, data = dfW.BCO))$coef
```

```
Estimate Std. Error  t value Pr(>|t|)
(Intercept)  0.40544    0.17976  2.25551 0.032412
Y.1          -0.38831    0.46152 -0.84139 0.407522
Y.3           0.36836    0.45811  0.80409 0.428370
```

but is quite close to the average period-specific ANCOVA treatment effect:

```
e.ANCOVA2 <- lm(Y.4 ~ Y.3 + treatment.4, data = dfW.BCO)
mean(c(summary(e.ANCOVA)$coef[3,"Estimate"],summary(e.ANCOVA2)$coef[3,"Estimate"]))
```

```
[1] 0.43672
```

⚠ In many applications, however, a treatment dependent variance and correlation is preferable:

```
eBCOT.lmm <- lmm(Y ~ period + sequence + treatment, repetition =~time|id,
               structure = CS(list(~treatment,~treatment)), data = dfL.BCO)
summary(anova(eBCOT.lmm, effects = c("treatmentB-treatmentA=0")))
```

Hypothesis-specific Wald test

```
              estimate      se   df lower upper p.value
treatmentB-treatmentA=0    0.414 0.172 29 0.062 0.765  0.0226 *
-----
:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1.
df: Satterthwaite approximation w.r.t. model-based se.
se: based on the observed information (model-based).
```

Appendix A Mixed model estimator as a weighted average

A.1 Treatment as repetition variable

Consider the linear mixed model matching the t-test results when estimating the treatment effect:

```
eTreat.lmm2tt <- lmm(duration ~ treatment, repetition = ~treatment|id, data = bloodpressureL)
coef(eTreat.lmm2tt)
```

```
(Intercept) treatmentB treatmentC
      1.7250      0.5750      1.2583
```

The estimates correspond to a Generalized Least Squared (GLS) estimator defined by:

- a block diagonal covariance matrix with element
- a design matrix with element:

```
Omega1 <- sigma(eTreat.lmm2tt,
                cluster = "all", simplify = TRUE)
Omega1[1:3,1:3]
```

```
      [,1]      [,2]      [,3]
[1,] 0.3347727 -0.0072727 0.047727
[2,] -0.0072727 0.1236364 0.162727
[3,] 0.0477273 0.1627273 0.372424
```

```
X1 <- model.matrix(eTreat.lmm2tt)
head(X1,3)
```

```
(Intercept) treatmentB treatmentC
1           1           0           0
2           1           1           0
3           1           0           1
```

The corresponding projector weight each observation:

- proportionally to the sample size for treatments related to the regression parameter
- 0 otherwise:

```
P1 <- solve(t(X1) %*% solve(Omega1) %*% X1) %*% t(X1) %*% solve(Omega1)
vecP1 <- apply(round(P1,4), MARGIN = 1, FUN = table, y = bloodpressureL$treatment)
vecP1
```

\$(Intercept)\$	\$treatmentB	\$treatmentC
A B C	A B C	A B C
0 0 12 12	-0.0833 12 0 0	-0.0833 12 0 0
0.0833 12 0 0	0 0 0 12	0 0 12 0
0.0833 12 0 0	0.0833 0 12 0	0.0833 0 0 12
total 1 0 0	total -1 1 0	total -1 0 1

We can verify that we retrieve the mixed model estimates:

```
(P1 %*% bloodpressureL$duration)[,1]
```

```
(Intercept) treatmentB treatmentC
      1.7250      0.5750      1.2583
```

A.2 Period as repetition variable

Consider the same linear mixed model but with period as repetition variable:

```
ePeriod.lmm2tt <- lmm(duration ~ treatment, repetition = ~period|id, data = bloodpressureL)
coef(ePeriod.lmm2tt)
```

```
(Intercept) treatmentB treatmentC
1.68755      0.58766      1.16557
```

The estimates correspond to a Generalized Least Squared (GLS) estimator defined by:

- a block diagonal covariance matrix with elements
- a design matrix with element:

```
Omega2 <- sigma(ePeriod.lmm2tt,
                cluster = "all", simplify = TRUE)
Omega2[1:3,1:3]
```

```
      [,1]      [,2]      [,3]
[1,] 0.229440 0.082455 0.01444
[2,] 0.082455 0.249826 0.11704
[3,] 0.014440 0.117040 0.36480
```

```
X2 <- model.matrix(eTreat.lmm2tt)
X2[1:3,1:3]
```

```
(Intercept) treatmentB treatmentC
1           1           0           0
2           1           1           0
3           1           0           1
```

The weighing of the observations is less intuitive as all treatments contribute, to various extends, to each regression parameter.

```
P2 <- solve(t(X2) %*% solve(Omega2) %*% X2) %*% t(X2) %*% solve(Omega2)
vecP2 <- apply(round(P2,4), MARGIN = 1, FUN = table, y = bloodpressureL$treatment,
               simplify = FALSE)
```

\$'(Intercept)'	\$treatmentB	\$treatmentC
A B C	A B C	A B C
-0.0156 0 4 0	-0.1052 4 0 0	-0.1078 4 0 0
-0.0129 0 0 4	-0.102 4 0 0	-0.0734 4 0 0
-0.0123 0 4 0	-0.0429 4 0 0	-0.0688 4 0 0
-0.0053 0 0 4	-0.0306 0 0 4	-0.0332 0 4 0
0.0182 0 0 4	-0.0026 0 0 4	0.0026 0 4 0
0.0279 0 4 0	0.0332 0 0 4	0.0306 0 4 0
0.0611 4 0 0	0.0688 0 4 0	0.0429 0 0 4
0.0922 4 0 0	0.0734 0 4 0	0.102 0 0 4
0.0967 4 0 0	0.1078 0 4 0	0.1052 0 0 4
total 1 0 0	total -1 1 0	total -1 0 1

We can verify that we retrieve the mixed model estimates:

```
(P2 %*% bloodpressureL$duration)[,1]
```

```
(Intercept) treatmentB treatmentC
1.68755      0.58766      1.16557
```

A.3 Treatment as repetition variable, adjusted for period

We now consider the linear mixed model similar to the t-test but adjusting for period:

```
eTreat.lmm <- lmm(duration ~ treatment + period, repetition = ~treatment|id,  
                  data = bloodpressureL)  
coef(eTreat.lmm)
```

```
(Intercept)  treatmentB  treatmentC    period2    period3  
    1.54915    0.57500    1.25833    0.19991    0.32764
```

As before we extract the residual-variance covariance matrix and the design matrix:

```
Omega3 <- sigma(eTreat.lmm,  
                cluster = "all", simplify = TRUE)  
X3 <- model.matrix(eTreat.lmm)
```

to understand how the GLS estimator weight each observation:

```
P3 <- solve(t(X3) %*% solve(Omega3) %*% X3) %*% t(X3) %*% solve(Omega3)  
vecP3 <- apply(round(P3,4), MARGIN = 1, FUN = table, y = bloodpressureL$treatment)  
vecP3
```

```
$ '(Intercept) '  
      A B C  
-0.0828 0 4 0  
-0.0497 0 0 4  
-0.0134 0 4 0  
-0.0048 0 0 4  
0.0546  0 0 4  
0.0631  4 0 0  
0.0876  4 0 0  
0.0962  0 4 0  
0.0992  4 0 0  
total   1 0 0  
  
$treatmentB  
      A B C  
-0.0833 12 0 0  
0        0 0 12  
0.0833  0 12 0  
total   -1 1 0  
  
$treatmentC  
      A B C  
-0.0833 12 0 0  
0        0 12 0  
0.0833  0 0 12  
total   -1 0 1
```

We can verify that we retrieve the mixed model estimates:

```
(P3 %*% bloodpressureL$duration)[,1]
```

```
(Intercept)  treatmentB  treatmentC    period2    period3  
    1.54915    0.57500    1.25833    0.19991    0.32764
```

A.4 Period as repetition variable, adjusted for period

We now consider the linear mixed model similar to the t-test but adjusting for period:

```
ePeriod.lmm <- lmm(duration ~ treatment + period, repetition = ~period|id,  
                  data = bloodpressureL)  
coef(ePeriod.lmm)
```

```
(Intercept)  treatmentB  treatmentC    period2    period3  
      1.31867      0.74657      1.39742      0.35833      0.55000
```

As before we extract the residual-variance covariance matrix and the design matrix:

```
Omega4 <- sigma(ePeriod.lmm,  
               cluster = "all", simplify = TRUE)  
X4 <- model.matrix(ePeriod.lmm)
```

to understand how the GLS estimator weight each observation:

```
P4 <- solve(t(X4) %*% solve(Omega4) %*% X4) %*% t(X4) %*% solve(Omega4)  
vecP4 <- apply(round(P4,4), MARGIN = 1, FUN = table, y = bloodpressureL$treatment)  
vecP4
```

\$(Intercept)\$		\$treatmentB		\$treatmentC
A B C		A B C		A B C
-0.0442 0 4 0		-0.1389 4 0 0		-0.1329 4 0 0
-0.0385 0 0 4		-0.0738 4 0 0		-0.0849 4 0 0
0.0032 0 4 0		-0.0477 0 0 4		-0.0417 0 4 0
0.0035 0 0 4		-0.0373 4 0 0		-0.0321 4 0 0
0.035 0 0 4		0.006 0 0 4		-0.006 0 4 0
0.0353 4 0 0		0.0321 0 4 0		0.0373 0 0 4
0.0407 4 0 0		0.0417 0 0 4		0.0477 0 4 0
0.041 0 4 0		0.0849 0 4 0		0.0738 0 0 4
0.1739 4 0 0		0.1329 0 4 0		0.1389 0 0 4
total 1 0 0		total -1 1 0		total -1 0 1

We can verify that we retrieve the mixed model estimates:

```
(P4 %*% bloodpressureL$duration)[,1]
```

```
(Intercept)  treatmentB  treatmentC    period2    period3  
      1.31867      0.74657      1.39742      0.35833      0.55000
```