Overview of the package LMMstar

Brice Ozenne

July 9, 2021

This vignette describes the main functionalities of the **LMMstar** package. This package implements specific types of linear mixed models mainly useful when having repeated observations over a discrete variable (e.g. time, brain region, . . .). Key assumptions are that at the cluster level, observation are independent and identically distributed and that the mean and variance are driven by independent factors. In particular, in large samples the residuals do not have to be normally distributed.

The **LMMstar** package contains four main functions:

- the function `lmm` is the main function of the package which fits linear mixed models. The user can interact with *lmm* objects using:

    - `anova` to test combinations of coefficients (Wald test or Likelihood ratio tests)
    - `coef` to extract the estimates.
    - `confint` to extract estimates, confidence intervals, and p.values.
    - `getVarCov` to extract the modeled residual variance covariance matrix.
    - `logLik` to output the log-likelihood of the estimated model.
    - `predict` to compute the conditional mean for new observations.
    - `residuals` to extract the observed residuals of the fitted model.
    - `summary` to obtain a summary of the results

- the `summarize` function to compute summary statistics stratified on a categorical variable (typically time).

- the `sampleRem` function to simulate longitudinal data.

- the `LMMstar.options` function enables the user to display the default values used in the **LMMstar** package. function. The function can also change the default values to better match the user needs.

Before going further we need to load the **LMMstar** package in the R session:

```
library(LMMstar)
```

To illustrate the functionalities of the package, we will use the `veteran` dataset:

```
data(gastricbypassL, package = "LMMstar")
head(gastricbypassL)
```

```
  id visit                    time weight glucagon
1 1      1 3 months before surgery  127.2  5032.50
2 2      1 3 months before surgery  165.2 12142.50
3 3      1 3 months before surgery  109.7 10321.35
4 4      1 3 months before surgery  146.2  6693.00
5 5      1 3 months before surgery  113.1  7090.50
6 6      1 3 months before surgery  158.8 10386.00
```

See `?gastricbypassL` for a presentation of the database. We will use a shorter version of the time variable:

```
gastricbypassL$time <- factor(gastricbypassL$time,
        levels = c("3 months before surgery", "1 week before surgery",
          "1 week after surgery", "3 months after surgery" ),
        labels = c("B3_months","B1_week","A1_week","A3_months"))
```

and rescale the glucagon values

```
gastricbypassL$glucagon <- as.double(scale(gastricbypassL$glucagon))
```

Note: the **LMMstar** package is under active development. Newer package versions may include additional functionalities and fix previous bugs. The version of the package that is being is:

```
utils::packageVersion("LMMstar")
```

```
[1] '0.2.3'
```

1 Descriptive statistics

Mean, standard deviation, and other summary statistic can be computed with respect to a categorical variable (typically time) using the `summarize` function:

```
sss <- summarize(weight+glucagon ~ time, data = gastricbypassL, na.rm = TRUE)
print(sss, digits = 3)
```

	outcome	time	observed	missing	mean	sd	min	median	max
1	weight	B3_months	20	0	128.9700	20.269	100.900	123.1000	173.000
2	weight	B1_week	20	0	121.2400	18.910	95.700	114.5000	162.200
3	weight	A1_week	20	0	115.7000	18.275	89.900	110.6000	155.000
4	weight	A3_months	20	0	102.3650	17.054	78.800	98.5000	148.000
5	glucagon	B3_months	20	0	-0.4856	0.641	-1.395	-0.6679	1.030
6	glucagon	B1_week	19	1	-0.6064	0.558	-1.416	-0.7669	0.946
7	glucagon	A1_week	19	1	1.0569	1.044	-0.478	0.9408	3.267
8	glucagon	A3_months	20	0	0.0576	0.760	-1.047	0.0319	2.124

2 Linear mixed model

2.1 Modeling tools

Fit a linear mixed model with **compound symmetry** structure:

```
eCS.lmm <- lmm(weight ~ time + glucagon,
        structure = CS(~time|id),
        data = gastricbypassL)
eCS.lmm
```

```
  Linear Mixed Model with a compound symmetry covariance matrix

data          : 78 observations and distributed in 20 clusters
log-likelihood : -243.6005
parameters    : 5 mean ((Intercept) timeB1_week timeA1_week timeA3_months glucagon)
                1 variance (sigma)
                1 correlation (Rho)
```

Fit a linear mixed model with **unstructured** covariance matrix:

```
eUN.lmm <- lmm(weight ~ time + glucagon,
        structure = UN(~time|id),
        data = gastricbypassL)
eUN.lmm
```

```
  Linear Mixed Model with an unstructured covariance matrix

data          : 78 observations and distributed in 20 clusters
log-likelihood : -216.3189
parameters    : 5 mean ((Intercept) timeB1_week timeA1_week timeA3_months glucagon)
                4 variance (sigma k.B1_week k.A1_week k.A3_months)
                6 correlation (cor(B1_week,B3_months) cor(A1_week,B3_months) cor(A3_months,B3_mont
```

Note: the calculation of the degrees of freedom, especially when using the observed information can be quite slow. Setting the arguments `df` to `FALSE` and `type.information` to `"expected"` when calling `lmm` should lead to a more reasonnable computation time.

2.2 Model output

The `summary` method can be used to display the main information relative to the model fit:

```
summary(eCS.lmm, ci = TRUE)
```

```
  Linear Mixed Model with a compound symmetry covariance matrix
  - fitted using Restricted Maximum Likelihood (REML)
  - log-likelihood :-243.6005 (parameters: mean = 5, variance = 1, correlation = 1)

Dataset: gastricbypassL
 - 20 clusters
 - 78 observations were analyzed, 2 were excluded because of missing values
 - 4 maximum number of observations per cluster
 - levels of the categorical variables
 - reference level: time=B3_months


$time
          B1_week A1_week A3_months
B3_months       0       0         0
B1_week         1       0         0
A1_week         0       1         0
A3_months       0       0         1


Correlation structure: ~1 | id
          B3_months B1_week A1_week A3_months
B3_months      1.00    0.97    0.97      0.97
B1_week        0.97    1.00    0.97      0.97
A1_week        0.97    0.97    1.00      0.97
A3_months      0.97    0.97    0.97      1.00


Variance structure: ~1
      standard.deviation
sigma           18.84957


Mean structure: weight ~ time + glucagon
              estimate    se     df   lower    upper p.value
(Intercept)    129.369 4.226 20.034 120.556  120.556  <0.001 ***
timeB1_week     -7.619 1.054 53.968  -9.732   -9.732  <0.001 ***
timeA1_week    -14.495 1.428 53.879 -17.358  -17.358  <0.001 ***
timeA3_months  -27.051 1.087 53.943 -29.231  -29.231  <0.001 ***
glucagon         0.822  0.62  53.81  -0.421   -0.421   0.191


The columns lower and upper correspond to the 95% confidence interval of the estimated coefficient
Note: p-values and confidence intervals are not adjusted for multiple comparisons
```

2.3 Extract estimated coefficients

The value of the estimated coefficients can be output using `coef`:

```
coef(eCS.lmm)
```

```
(Intercept)    timeB1_week    timeA1_week  timeA3_months      glucagon         sigma           Rho
 129.3690995     -7.6194918    -14.4951323    -27.0514694     0.8217879    18.8495684     0.9699341
```

It is possible to apply specific transformation on the variance coefficients, for instance to obtain the residual variance relative to each outcome:

```
coef(eUN.lmm, effects = "variance", transform.k = "sd")
```

```
sigma:B3_months    sigma:B1_week    sigma:A1_week  sigma:A3_months
       20.28080         19.04553         17.65479         16.76104
```

2.4 Extract estimated residual variance-covariance structure

The method `getVarCov` can be used to output the covariance structure of the residuals:

```
nlme::getVarCov(eCS.lmm)
```

```
          B3_months  B1_week  A1_week A3_months
B3_months  355.3062 344.6236 344.6236  344.6236
B1_week    344.6236 355.3062 344.6236  344.6236
A1_week    344.6236 344.6236 355.3062  344.6236
A3_months  344.6236 344.6236 344.6236  355.3062
```

It can also be specific to an individual:

```
nlme::getVarCov(eCS.lmm, individual = 5)
```

```
          B3_months  A1_week A3_months
B3_months  355.3062 344.6236  344.6236
A1_week    344.6236 355.3062  344.6236
A3_months  344.6236 344.6236  355.3062
```

2.5 Model diagnostic

The method `residuals` can also be used to extract the residulas in the wide format:

```
eCS.diagW <- residuals(eCS.lmm, type.residual = "normalized", format = "wide")
head(eCS.diagW)
```

```
  cluster  B3_months        B1_week      A1_week   A3_months
1       1 -0.8042448 -0.709908611 -1.4242831   0.3176640
2       2  1.0863177 -0.133256804  1.1083627   1.5977042
3       3 -0.4597852 -0.612727870 -0.6060136  -0.8589524
4       4 -1.0103075  0.007471088  0.1309862   1.1428822
5       5 -0.1258773           NA -0.3819184  -0.7874832
6       6  3.5646225  2.333205076  2.8387204   0.3586263
```

or in the long format:

```
eCS.diagL <- residuals(eCS.lmm, type.residual = "normalized", format = "long")
head(eCS.diagL)
```

```
[1] -0.8042448  1.0863177 -0.4597852 -1.0103075 -0.1258773  3.5646225
```

Various type of residuals can be extract but the normalized one are recommanded when doing model checking. The method `residuals` can also be used to display diagnostic plots, e.g. about:

- the distribution of the residuals across fitted values using a scatterplot

```
residuals(eCS.lmm, type.residual = "normalized", plot = "scatterplot", size.text = 20)
```

- the "normality" of the residuals at each repetition using a quantile-quantile plot [1]:

```
residuals(eCS.lmm, type.residual = "normalized", format = "wide",
    plot = "qqplot", engine.qqplot = "qqtest")
## Note: the qqtest package to be installed to use the argument engine.plot = "qqtest"
```
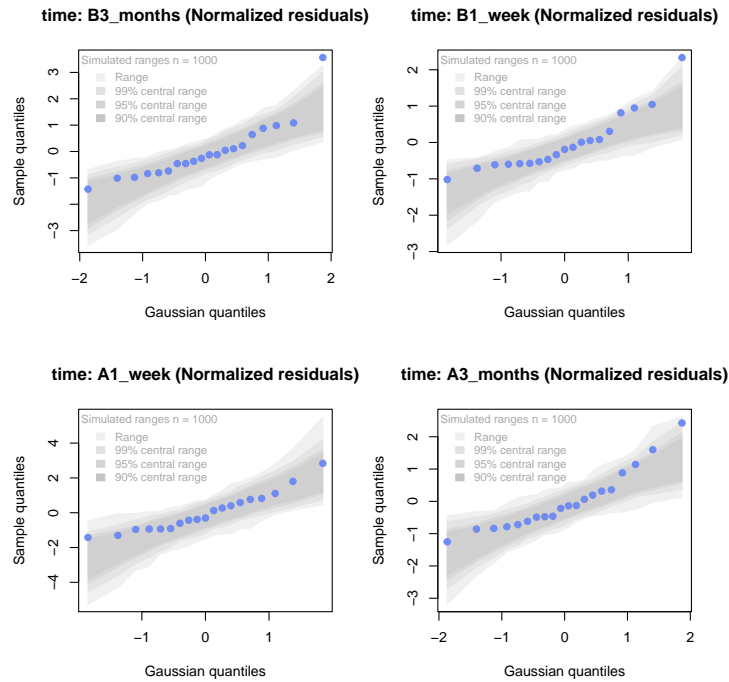

- the residual correlation within cluster between the residuals:

```
residuals(eCS.lmm, type.residual = "normalized", plot = "correlation", format = "wide",
    size.text = 20)
```


Normalized residuals

---
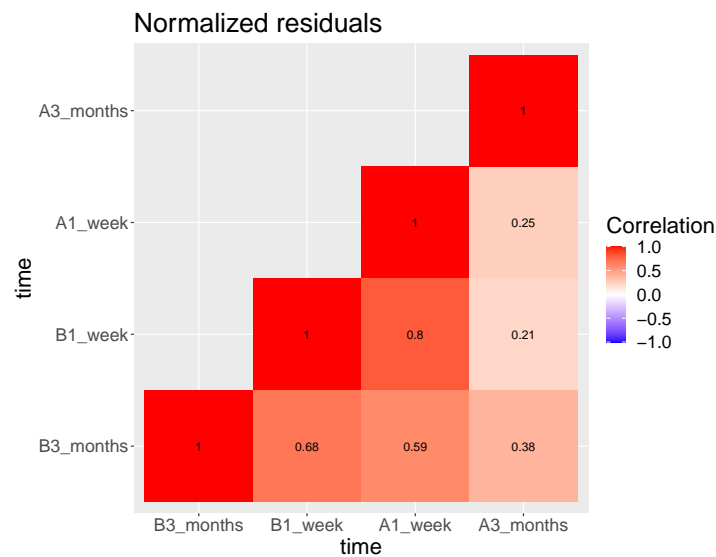
[1]see Oldford (2016) for guidance about how to read quantile-quantile plots.

2.6 Model fit

The fitted values can be displayed via the `emmeans` package or using the `autoplot` method:

```
library(emmeans) ## left panel
emmip(eCS.lmm, ~time) + theme(text = element_text(size=20))
```

```
library(ggplot2) ## right panel
autoplot(eCS.lmm, color = "id", size.text = 20)
```


In the first case the average curve (over glucago values) is displayed while in the latter each possible curve is displayed. With the `autoplot` method, it is possible to display a curve specific to a glucagon value via the argument `at`:

```
autoplot(eCS.lmm, at = data.frame(glucagon = 10), color = "glucagon")
```

2.7 Statistical inference

2.7.1 Model coefficients

The estimated coefficients with their confidence intervals can be accessed via the `confint` method:

```
confint(eCS.lmm)
```

```
                estimate        se       df      lower        upper      p.value
(Intercept)    129.3690995 4.2256318 20.03432 120.5555539 138.1826451 0.000000e+00
timeB1_week     -7.6194918 1.0538287 53.96824  -9.7323197  -5.5066639 1.746652e-09
timeA1_week    -14.4951323 1.4279524 53.87927 -17.3581515 -11.6321131 4.130030e-14
timeA3_months  -27.0514694 1.0870651 53.94292 -29.2309565 -24.8719823 0.000000e+00
glucagon         0.8217879 0.6199685 53.80984  -0.4212748   2.0648506 1.905952e-01
sigma           18.8495684 0.1587900 16.79510  13.4793760   26.3592491          NA
Rho              0.9699341 0.1874830 29.43362   0.9363994    0.9859156 4.330536e-12
Note: estimates and confidence intervals for sigma, rho have been back-transformed.
     standard errors are not back-transformed.
```

The variance and correlation parameters being constrained parameters (e.g. strictly positive), they uncertainty is by default computed after transformation (e.g. `log`) and then backtransformed. The transformed values can be displayed using:

```
confint(eCS.lmm, backtransform = FALSE)
```

```
                estimate        se       df       lower       upper      p.value
(Intercept)    129.3690995 4.2256318 20.03432 120.5555539 138.182645 0.000000e+00
timeB1_week     -7.6194918 1.0538287 53.96824  -9.7323197  -5.506664 1.746652e-09
timeA1_week    -14.4951323 1.4279524 53.87927 -17.3581515 -11.632113 4.130030e-14
timeA3_months  -27.0514694 1.0870651 53.94292 -29.2309565 -24.871982 0.000000e+00
glucagon         0.8217879 0.6199685 53.80984  -0.4212748   2.064851 1.905952e-01
log(sigma)       2.9364900 0.1587900 16.79510   2.6011608   3.271819           NA
atanh(Rho)       2.0911816 0.1874830 29.43362   1.7079810   2.474382 4.330536e-12
```

2.7.2   Linear combination of the model coefficients

The `anova` method can be use to test one or several linear combinations of the model coefficients using Wald tests. For instance whether there is a change in average weight just after taking the treatment:

```
anova(eUN.lmm, effects = c("timeA1_week-timeB1_week=0"), ci = TRUE)
```

```
                   ** User-specified hypotheses **
 - F-test
 statistic df.num df.denom      p.value
  43.14145      1 17.87461 3.723244e-06


 - P-values and confidence interval (adjusted for multiplicity within each global test)
                            estimate        se       df     lower     upper      p.value
timeA1_week - timeB1_week -3.905721 0.5946396 17.87461 -5.155641 -2.655801 3.723244e-06
```

When testing transformed variance or correlation parameters, parentheses (as in `log(k).B1_week`) cause problem for recognizing parameters:

```
try(
  anova(eUN.lmm,
 effects = c("log(k).B1_week=0","log(k).A1_week=0","log(k).A3_months=0"))
)
```

```
Error in .anova_Wald(object, effects = effects, rhs = rhs, df = df, ci = ci,  :
  Possible mispecification of the argument 'effects' as running mulcomp::glht lead to the following
Error in parse(text = ex[i]) : <text>:1:7: unexpected symbol
1: log(k).B1_week
          ^
```

It is then advised to specify the null hypothesis via a contrast matrix, e.g.:

```
name.coef <- rownames(confint(eUN.lmm, backtransform = FALSE))
name.varcoef <- grep("log(k)",name.coef, value = TRUE, fixed = TRUE)
C <- matrix(0, nrow = 3, ncol = length(name.coef), dimnames = list(name.varcoef, name.coef))
diag(C[name.varcoef,name.varcoef]) <- 1

anova(eUN.lmm, effects = C)
```

```
                  ** User-specified hypotheses **
- F-test
statistic df.num df.denom      p.value
 6.203176      3 17.99457 0.004417067
```

2.8 Baseline adjustment

The `lmm` contains an "experimental" feature to drop non-identifiable effects from the model. For instance, let us define two (artifical) groups of patients:

```
gastricbypassL$group <- c("1","2")[as.numeric(gastricbypassL$id) %in% 15:20 + 1]
```

We would like to model group differences only after baseline (i.e. only at 1 week and 3 months after). For this we will define a treatment variable being the group variable except before baseline where it is "none":

```
gastricbypassL$treat <- baselineAdjustment(gastricbypassL, variable = "group",
        repetition = ~time|id, constrain = c("B3_months","B1_week"),
        new.level = "none")
table(treat = gastricbypassL$treat, time = gastricbypassL$time, group = gastricbypassL$group)
```

```
, , group = 1

      time
treat  B3_months B1_week A1_week A3_months
  none        14      14       0         0
  1            0       0      14        14
  2            0       0       0         0

, , group = 2

      time
treat  B3_months B1_week A1_week A3_months
  none         6       6       0         0
  1            0       0       0         0
  2            0       0       6         6
```

Here we will be able to estimate a total of 6 means and therefore can at most identify 6 effects. However the design matrix for the interaction model:

```
colnames(model.matrix(weight ~ treat*time, data = gastricbypassL))
```

```
[1] "(Intercept)"        "treat1"            "treat2"             "timeB1_week"
[5] "timeA1_week"        "timeA3_months"     "treat1:timeB1_week" "treat2:timeB1_week"
[9] "treat1:timeA1_week" "treat2:timeA1_week" "treat1:timeA3_months" "treat2:timeA3_months"
```

contains 12 parameters (i.e. 6 too many). The `lmm` function will internally remove the one that cannot be identified and fit a simplified model:

```
eC.lmm <- lmm(weight ~ treat*time, data = gastricbypassL, structure = UN(~time|id))
```

```
Warning message:
In model.matrix_regularize(formula.mean, data) :
  Constant values in the design matrix in interactions "treat:time"
 Coefficients "treat1" "treat2" "timeA1_week" "timeA3_months" "treat1:timeB1_week" "treat2:timeB1_w
Consider defining manually the interaction, e.g. via droplevels(interaction(.,.)) to avoid this war
```

with the following coefficients:

```
coef(eC.lmm, effects = "mean")
```

```
          (Intercept)              timeB1_week    treat1:timeA1_week    treat2:timeA1_week
            128.97000                 -7.73000             -12.83949             -14.27452
treat1:timeA3_months treat2:timeA3_months
            -27.07620                -25.50553
```

One can vizualize the baseline adjustment via the `autoplot` function:

```
autoplot(eC.lmm, color = "group", ci = FALSE, size.text = 20)
```


To more easily compare the two groups, one could set the baseline treatment to the treatment in the control arm by omitting the argument `new.level`:

```
gastricbypassL$treat2 <- baselineAdjustment(gastricbypassL, variable = "group",
        repetition = ~time|id, constrain = c("B3_months","B1_week"))
table(treat = gastricbypassL$treat2, time = gastricbypassL$time, group = gastricbypassL$group)
```

```
, , group = 1

     time
treat B3_months B1_week A1_week A3_months
    1         14      14      14        14
    2          0       0       0         0

, , group = 2

     time
treat B3_months B1_week A1_week A3_months
    1          6       6       0         0
    2          0       0       6         6
```

Fitting the model

```
eC2.lmm <- suppressWarnings(lmm(weight ~ treat2*time, data = gastricbypassL, structure = UN(~
   time|id)))
```

will directly output group differences:

```
confint(eC2.lmm, effects = "mean")[5:6,]
```

```
                     estimate        se        df      lower      upper    p.value
treat22:timeA1_week   -1.435033 0.621046 16.25085 -2.749942 -0.1201245 0.0342893
treat22:timeA3_months  1.570675 2.462555 16.28753 -3.642229  6.7835794 0.5324540
```

2.9   Marginal means

The `lmm` function can be used in conjonction with the `emmeans` package to compute marginal means. Consider the following model:

```
e.group <- lmm(weight ~ time*group, data = gastricbypassL, structure = UN(~time|id))
```

We can for instance compute the average value over time *assuming balanced groups*:

```
library(emmeans)
emmeans(e.group, specs=~time)
```

```
NOTE: Results may be misleading due to involvement in interactions
 time       emmean   SE   df lower.CL upper.CL
 B3_months    130 5.05 18.0    119.3      141
 B1_week      122 4.69 18.0    112.5      132
 A1_week      117 4.55 18.0    107.0      126
 A3_months    104 4.20 18.1     94.9      113

Results are averaged over the levels of: group
Confidence level used: 0.95
```

This differs from the average value over time over the whole sample:

```
df.pred <- cbind(gastricbypassL, predict(e.group, newdata = gastricbypassL))
summarize(formula = estimate~time, data = df.pred)
```

```
   outcome       time observed missing    mean       sd      min   median     max
1 estimate B3_months       20       0 128.970 2.270212 127.5214 127.5214 132.35
2 estimate   B1_week       20       0 121.240 2.726942 119.5000 119.5000 125.30
3 estimate   A1_week       20       0 115.700 2.014981 114.4143 114.4143 118.70
4 estimate A3_months       20       0 102.365 3.146729 100.3571 100.3571 107.05
```

as the groups are not balanced and with this approach more "weight" is given to the expected value group 1 as it contains more indiviuals.

```
table(group = gastricbypassL$group, time = gastricbypassL$time)
```

```
      time
group B3_months B1_week A1_week A3_months
    1        14      14      14        14
    2         6       6       6         6
```

By hand:

```
mu.group1 <-  as.double(coef(e.group)["(Intercept)"])
mu.group2 <-  as.double(coef(e.group)["(Intercept)"] + coef(e.group)["group2"])
p.group1 <- 14/20
p.group2 <- 6/20
c(emmeans = (mu.group1+mu.group2)/2,
  predict = mu.group1 * p.group1 + mu.group2 * p.group2)
```

```
 emmeans   predict
129.9357 128.9700
```

which one is relevant depends on the application. The `emmeans` function can also be used to display expected value in each group over time:

```
emmeans.group <- emmeans(e.group, specs = ~group|time)
emmeans.group
```

```
time = B3_months:
 group emmean   SE   df lower.CL upper.CL
 1        128 5.53 18.0    115.9      139
 2        132 8.45 18.0    114.6      150

time = B1_week:
 group emmean   SE   df lower.CL upper.CL
 1        120 5.14 18.0    108.7      130
 2        125 7.85 18.0    108.8      142

time = A1_week:
 group emmean   SE   df lower.CL upper.CL
 1        114 4.99 18.0    103.9      125
 2        119 7.62 18.0    102.7      135

time = A3_months:
 group emmean   SE   df lower.CL upper.CL
 1        100 4.60 18.1     90.7      110
 2        107 7.03 18.1     92.3      122

Confidence level used: 0.95
```

Using the `pair` function displays the differences:

```
epairs.group <- pairs(emmeans.group, reverse = TRUE)
epairs.group
```

```
time = B3_months:
 contrast estimate    SE   df t.ratio p.value
 2 - 1        4.83 10.10 18.0 0.478    0.6383


time = B1_week:
 contrast estimate    SE   df t.ratio p.value
 2 - 1        5.80  9.38 18.0 0.618    0.5441


time = A1_week:
 contrast estimate    SE   df t.ratio p.value
 2 - 1        4.29  9.11 18.0 0.471    0.6435


time = A3_months:
 contrast estimate    SE   df t.ratio p.value
 2 - 1        6.69  8.40 18.1 0.797    0.4361
```

One can adjust for multiple comparison via the `adjust` argument and display confidence intervals setting the argument `infer` to `TRUE`:

```
summary(epairs.group, by = NULL, adjust = "mvt", infer = TRUE)
```

```
 contrast time        estimate    SE   df lower.CL upper.CL t.ratio p.value
 2 - 1    B3_months       4.83 10.10 18.0    -18.0     27.6 0.478    0.7496
 2 - 1    B1_week         5.80  9.38 18.0    -15.4     27.0 0.618    0.6482
 2 - 1    A1_week         4.29  9.11 18.0    -16.3     24.9 0.471    0.7556
 2 - 1    A3_months       6.69  8.40 18.1    -12.3     25.7 0.797    0.5287


Confidence level used: 0.95
Conf-level adjustment: mvt method for 4 estimates
P value adjustment: mvt method for 4 tests
```

This should also work when doing baseline adjustment (because of baseline adjustment no difference is expected at the first two timepoints):

```
summary(pairs(emmeans(eC2.lmm , specs = ~treat2|time), reverse = TRUE), by = NULL)
```

```
contrast time        estimate    SE   df t.ratio p.value
2 - 1    B3_months       0.00 0.000  NaN    NaN     NaN
2 - 1    B1_week         0.00 0.000  NaN    NaN     NaN
2 - 1    A1_week        -1.44 0.621 16.2 -2.311  0.0345
2 - 1    A3_months       1.57 2.463 16.3  0.638  0.5326
```

3 Data generation

Simulate some data in the wide format:

```r
set.seed(10) ## ensure reproductibility
n.obs <- 100
n.times <- 4
mu <- rep(0,4)
gamma <- matrix(0, nrow = n.times, ncol = 10) ## add interaction
gamma[,6] <- c(0,1,1.5,1.5)
dW <- sampleRem(n.obs, n.times = n.times, mu = mu, gamma = gamma, format = "wide")
head(round(dW,3))
```

	id	X1	X2	X3	X4	X5	X6	X7	X8	X9	X10	Y1	Y2	Y3	Y4
1	1	1	0	1	1	0	-0.367	1.534	-1.894	1.729	0.959	1.791	2.429	3.958	2.991
2	2	1	0	1	2	0	-0.410	2.065	1.766	0.761	-0.563	2.500	4.272	3.002	2.019
3	3	0	0	2	1	0	-1.720	-0.178	2.357	1.966	1.215	-3.208	-5.908	-4.277	-5.154
4	4	0	0	0	1	0	0.923	-2.089	0.233	1.307	-0.906	-2.062	0.397	1.757	-1.380
5	5	0	0	2	1	0	0.987	5.880	0.385	0.028	0.820	7.963	7.870	7.388	8.609
6	6	0	0	1	1	2	-1.075	0.479	2.202	0.900	-0.739	0.109	-1.602	-1.496	-1.841

Simulate some data in the long format:

```r
set.seed(10) ## ensure reproductibility
dL <- sampleRem(n.obs, n.times = n.times, mu = mu, gamma = gamma, format = "long")
head(dL)
```

	id	visit	Y	X1	X2	X3	X4	X5	X6	X7	X8	X9	X10
1	1	1	1.791444	1	0	1	1	0	-0.3665251	1.533815	-1.894425	1.7288665	0.9592499
2	1	2	2.428570	1	0	1	1	0	-0.3665251	1.533815	-1.894425	1.7288665	0.9592499
3	1	3	3.958350	1	0	1	1	0	-0.3665251	1.533815	-1.894425	1.7288665	0.9592499
4	1	4	2.991198	1	0	1	1	0	-0.3665251	1.533815	-1.894425	1.7288665	0.9592499
5	2	1	2.500179	1	0	1	2	0	-0.4097541	2.065413	1.765841	0.7613348	-0.5630173
6	2	2	4.272357	1	0	1	2	0	-0.4097541	2.065413	1.765841	0.7613348	-0.5630173

4 Modifying default options

The `LMMstar.options` method enable to get and set the default options used by the package. For instance, the default option for the information matrix is:

```
LMMstar.options("type.information")
```

```
$type.information
[1] "observed"
```

To change the default option to "expected" (faster to compute but less accurate p-values and confidence intervals in small samples) use:

```
LMMstar.options(type.information = "expected")
```

To restore the original default options do:

```
LMMstar.options(reinitialise = TRUE)
```

5   R session

Details of the R session used to generate this document:

```
sessionInfo()
```

```
R version 4.1.0 (2021-05-18)
Platform: x86_64-pc-linux-gnu (64-bit)
Running under: Ubuntu 20.04.2 LTS

Matrix products: default
BLAS:   /usr/lib/x86_64-linux-gnu/blas/libblas.so.3.9.0
LAPACK: /usr/lib/x86_64-linux-gnu/lapack/liblapack.so.3.9.0

locale:
 [1] LC_CTYPE=en_US.UTF-8       LC_NUMERIC=C               LC_TIME=en_US.UTF-8
 [4] LC_COLLATE=en_US.UTF-8     LC_MONETARY=en_US.UTF-8    LC_MESSAGES=en_US.UTF-8
 [7] LC_PAPER=en_US.UTF-8       LC_NAME=C                  LC_ADDRESS=C
[10] LC_TELEPHONE=C             LC_MEASUREMENT=en_US.UTF-8 LC_IDENTIFICATION=C

attached base packages:
[1] stats     graphics  grDevices utils     datasets  methods   base

other attached packages:
 [1] sandwich_3.0-1    reshape2_1.4.4     numDeriv_2016.8-1.1 nlme_3.1-152
 [5] multcomp_1.4-17   TH.data_1.0-10     MASS_7.3-54          survival_3.2-11
 [9] mvtnorm_1.1-2     Matrix_1.3-4       lava_1.6.9           psych_2.1.3
[13] lattice_0.20-44   emmeans_1.6.0      qqtest_1.2.0         ggplot2_3.3.4
[17] LMMstar_0.2.3

loaded via a namespace (and not attached):
 [1] Rcpp_1.0.6        prettyunits_1.1.1 ps_1.6.0          zoo_1.8-9         rprojroot_2.0.2
 [6] digest_0.6.27     utf8_1.2.1        R6_2.5.0          plyr_1.8.6        coda_0.19-4
[11] pillar_1.6.1      rlang_0.4.11      data.table_1.14.0 callr_3.7.0       desc_1.3.0
[16] labeling_0.4.2    devtools_2.4.1    splines_4.1.0     stringr_1.4.0     munsell_0.5.0
[21] compiler_4.1.0    pkgconfig_2.0.3   mnormt_2.0.2      pkgbuild_1.2.0    tmvnsim_1.0-2
[26] mgcv_1.8-36       tidyselect_1.1.1  tibble_3.1.2      codetools_0.2-18  fansi_0.5.0
[31] crayon_1.4.1      dplyr_1.0.6       withr_2.4.2       grid_4.1.0        butils.base_1.2
[36] xtable_1.8-4      gtable_0.3.0      lifecycle_1.0.0   magrittr_2.0.1    scales_1.1.1
[41] pbapply_1.4-3     estimability_1.3  cli_2.5.0         stringi_1.6.2     cachem_1.0.5
[46] farver_2.1.0      remotes_2.3.0     fs_1.5.0          testthat_3.0.2    ellipsis_0.3.2
[51] generics_0.1.0    vctrs_0.3.8       tools_4.1.0       glue_1.4.2        purrr_0.3.4
[56] pkgload_1.2.1     processx_3.5.2    parallel_4.1.0    fastmap_1.1.0     colorspace_2.0-1
[61] sessioninfo_1.1.1 memoise_2.0.0     usethis_2.0.1
```

References

Oldford, R. W. (2016). Self-calibrating quantile–quantile plots. *The American Statistician*, 70(1):74–90.

Appendix A Likelihood in a linear mixed model

A.1 Log-likelihood

Denote by $\boldsymbol{Y}$ a vector of m outcomes, $\boldsymbol{X}$ a vector of p covariates, $\mu(\boldsymbol{\Theta}, \boldsymbol{X})$ the modeled mean, and $\Omega(\boldsymbol{\Theta}, \boldsymbol{X})$ the modeled residual variance-covariance. The restricted log-likelihood in a linear mixed model can then be written:

$$
\begin{aligned}
\mathcal{L}(\boldsymbol{\Theta}|\boldsymbol{Y}, \boldsymbol{X}) = &\frac{p}{2}\log(2\pi) - \frac{1}{2}\log\left(\left|\sum_{i=1}^{n} \boldsymbol{X}_i \Omega_i^{-1}(\boldsymbol{\Theta})\boldsymbol{X}_i^{\intercal}\right|\right) \\
&+ \sum_{i=1}^{n}\left(-\frac{m}{2}\log(2\pi) - \frac{1}{2}\log|\Omega_i(\boldsymbol{\Theta})| - \frac{1}{2}(\boldsymbol{Y}_i - \mu(\boldsymbol{\Theta}, \boldsymbol{X}_i))\Omega_i(\boldsymbol{\Theta})^{-1}(\boldsymbol{Y}_i - \mu(\boldsymbol{\Theta}, \boldsymbol{X}_i))^{\intercal}\right) \quad \text{(A)}
\end{aligned}
$$

This is what the `logLik` method is computing for the REML criteria. The red term is specific to the REML criteria and prevents from computing individual contributions to the likelihood[2]. The blue term is what `logLik` outputs for the ML criteria when setting the argument `indiv` to `TRUE`.

A.2 Score

Using that $\partial \log(\det(X)) = tr(X^{-1}\partial(X))$, the score is obtained by derivating once the log-likelihood, i.e., for $\theta \in \boldsymbol{\Theta}$:

$$
\begin{aligned}
\mathcal{S}(\theta) = &\frac{\partial \mathcal{L}(\boldsymbol{\Theta}|\boldsymbol{Y}, \boldsymbol{X})}{\partial \theta} = \frac{1}{2}tr\left(\left(\sum_{i=1}^{n} \boldsymbol{X}_i \Omega_i^{-1}(\boldsymbol{\Theta})\boldsymbol{X}_i^{\intercal}\right)^{-1}\left(\sum_{i=1}^{n} \boldsymbol{X}_i \Omega_i^{-1}(\boldsymbol{\Theta})\frac{\partial \Omega_i(\boldsymbol{\Theta})}{\partial \theta}\Omega_i(\boldsymbol{\Theta})^{-1}\boldsymbol{X}_i^{\intercal}\right)\right) \\
&+ \sum_{i=1}^{n}\left(-\frac{1}{2}tr\left(\Omega_i(\boldsymbol{\Theta})^{-1}\frac{\partial \Omega_i(\boldsymbol{\Theta})}{\partial \theta}\right) + \frac{\partial \mu(\boldsymbol{\Theta}, \boldsymbol{X}_i)}{\partial \theta}\Omega_i(\boldsymbol{\Theta})^{-1}(\boldsymbol{Y}_i - \mu(\boldsymbol{\Theta}, \boldsymbol{X}_i))^{\intercal}\right. \\
&\left.+ \frac{1}{2}(\boldsymbol{Y}_i - \mu(\boldsymbol{\Theta}, \boldsymbol{X}_i))\Omega_i(\boldsymbol{\Theta})^{-1}\frac{\partial \Omega_i(\boldsymbol{\Theta})}{\partial \theta}\Omega_i(\boldsymbol{\Theta})^{-1}(\boldsymbol{Y}_i - \mu(\boldsymbol{\Theta}, \boldsymbol{X}_i))^{\intercal}\right).
\end{aligned}
$$

This is what the `score` method is computing for the REML criteria. The red term is specific to the REML criteria and prevents from computing the score relative to each cluster. The blue term is what `score` outputs for the ML criteria when setting the argument `indiv` to `TRUE`.

---

[2]The REML is the likelihood of the observations divided by the prior on the estimated mean parameters $\widehat{\boldsymbol{\Theta}}_\mu \sim \mathcal{N}(\mu, (\boldsymbol{X}\Omega^{-1}(\boldsymbol{\Theta})\boldsymbol{X}^{\intercal})^{-1})$. This corresponds to $\frac{1}{\sqrt{2\pi}^p|(\sum_{i=1}^{n} \boldsymbol{X}_i \Omega_i^{-1}(\boldsymbol{\Theta})\boldsymbol{X}_i^{\intercal})^{-1}|}\exp\left(-(\widehat{\boldsymbol{\Theta}}_\mu - \mu)\left(2\sum_{i=1}^{n} \boldsymbol{X}_i \Omega_i^{-1}(\boldsymbol{\Theta})\boldsymbol{X}_i^{\intercal}\right)^{-1})(\widehat{\boldsymbol{\Theta}}_\mu - \mu)^{\intercal}\right)$ Since μ will be estimated to be $\boldsymbol{\Theta}_\mu$, the exponential term equals 1 and thus does not contribute to the log-likelihood. One divided by the other term gives $\sqrt{2\pi}^p\left(|\sum_{i=1}^{n} \boldsymbol{X}_i \Omega_i^{-1}(\boldsymbol{\Theta})\boldsymbol{X}_i^{\intercal}|\right)^{-1}$. The log of this term equals the red term

A.3 Hessian

Derivating a second time the log-likelihood gives the hessian, $\mathcal{H}(\boldsymbol{\Theta})$, with element[3]:

$$\mathcal{H}(\theta, \theta') = \frac{\partial^2 \mathcal{L}(\boldsymbol{\Theta}|\boldsymbol{Y}, \boldsymbol{X})}{\partial\theta\partial\theta'} = \frac{\partial \mathcal{S}(\theta)}{\partial\theta'}$$

$$= \frac{1}{2}tr\left(\left(\sum_{i=1}^{n} \boldsymbol{X}_i\Omega_i^{-1}(\boldsymbol{\Theta})\boldsymbol{X}_i^{\intercal}\right)^{-1}\left\{\sum_{i=1}^{n} \boldsymbol{X}_i\Omega_i^{-1}(\boldsymbol{\Theta})\left(\frac{\partial^2\Omega_i(\boldsymbol{\Theta})}{\partial\theta\partial\theta'} - 2\frac{\partial\Omega_i(\boldsymbol{\Theta})}{\partial\theta}\Omega_i^{-1}(\boldsymbol{\Theta})\frac{\partial\Omega_i(\boldsymbol{\Theta})}{\partial\theta'}\right)\Omega_i(\boldsymbol{\Theta})^{-1}\boldsymbol{X}_i^{\intercal}\right.$$

$$\left.+ \left(\sum_{i=1}^{n} \boldsymbol{X}_i\Omega_i^{-1}(\boldsymbol{\Theta})\frac{\partial\Omega_i(\boldsymbol{\Theta})}{\partial\theta}\Omega_i(\boldsymbol{\Theta})^{-1}\boldsymbol{X}_i^{\intercal}\right)\left(\sum_{i=1}^{n} \boldsymbol{X}_i\Omega_i^{-1}(\boldsymbol{\Theta})\boldsymbol{X}_i^{\intercal}\right)^{-1}\left(\sum_{i=1}^{n} \boldsymbol{X}_i\Omega_i^{-1}(\boldsymbol{\Theta})\frac{\partial\Omega_i(\boldsymbol{\Theta})}{\partial\theta'}\Omega_i(\boldsymbol{\Theta})^{-1}\boldsymbol{X}_i^{\intercal}\right)\right\}\right)$$

$$+ \sum_{i=1}^{n}\left(\frac{1}{2}tr\left(\Omega_i(\boldsymbol{\Theta})^{-1}\frac{\partial\Omega_i(\boldsymbol{\Theta})}{\partial\theta'}\Omega_i(\boldsymbol{\Theta})^{-1}\frac{\partial\Omega_i(\boldsymbol{\Theta})}{\partial\theta} - \Omega_i(\boldsymbol{\Theta})^{-1}\frac{\partial^2\Omega_i(\boldsymbol{\Theta})}{\partial\theta\partial\theta'}\right)\right.$$

$$- \frac{\partial\mu(\boldsymbol{\Theta}, \boldsymbol{X}_i)}{\partial\theta}\Omega_i(\boldsymbol{\Theta})^{-1}\frac{\partial\Omega_i(\boldsymbol{\Theta})^{-1}}{\partial\theta'}\Omega_i(\boldsymbol{\Theta})^{-1}\boldsymbol{\varepsilon}_i(\boldsymbol{\Theta})^{\intercal} - \frac{\partial\mu(\boldsymbol{\Theta}, \boldsymbol{X}_i)}{\partial\theta}\Omega_i(\boldsymbol{\Theta})^{-1}\frac{\partial\mu(\boldsymbol{\Theta}, \boldsymbol{X}_i)}{\partial\theta'}^{\intercal}$$

$$\left.+ \frac{1}{2}\boldsymbol{\varepsilon}_i(\boldsymbol{\Theta})\Omega_i(\boldsymbol{\Theta})^{-1}\left(\frac{\partial^2\Omega_i(\boldsymbol{\Theta})}{\partial\theta\partial\theta'} - \frac{\partial\Omega_i(\boldsymbol{\Theta})}{\partial\theta'}\Omega_i(\boldsymbol{\Theta})^{-1}\frac{\partial\Omega_i(\boldsymbol{\Theta})}{\partial\theta} - \frac{\partial\Omega_i(\boldsymbol{\Theta})}{\partial\theta}\Omega_i(\boldsymbol{\Theta})^{-1}\frac{\partial\Omega_i(\boldsymbol{\Theta})}{\partial\theta'}\right)\Omega_i(\boldsymbol{\Theta})^{-1}\boldsymbol{\varepsilon}_i(\boldsymbol{\Theta})^{\intercal}\right).$$

where $\boldsymbol{\varepsilon}_i(\boldsymbol{\Theta}) = \boldsymbol{Y}_i - \mu(\boldsymbol{\Theta}, \boldsymbol{X}_i)$.

The `information` method will (by default) return the (observed) information which is the opposite of the hessian. So multiplying the previous formula by -1 gives what `information` output for the REML criteria. The red term is specific to the REML criteria and prevents from computing the information relative to each cluster. The blue term is what `information` outputs for the ML criteria (up to a factor -1) when setting the argument `indiv` to `TRUE`.

A possible simplification is to use the expected hessian. Indeed for any deterministic matrix A:

- $\mathbb{E}\left[A(\boldsymbol{Y}_i - \mu(\boldsymbol{\Theta}, \boldsymbol{X}_i))^{\intercal}|\boldsymbol{X}_i\right] = 0$

- $\mathbb{E}\left[(\boldsymbol{Y}_i - \mu(\boldsymbol{\Theta}, \boldsymbol{X}_i))A(\boldsymbol{Y}_i - \mu(\boldsymbol{\Theta}, \boldsymbol{X}_i))^{\intercal}||\boldsymbol{X}_i\right] = tr(A\mathbb{V}ar(\boldsymbol{Y}_i - \mu(\boldsymbol{\Theta}, \boldsymbol{X}_i)))$

Leading to:

$$\mathbb{E}\left[\mathcal{H}(\theta, \theta')|\boldsymbol{X}\right]$$

$$= \frac{1}{2}tr\left(\left(\sum_{i=1}^{n} \boldsymbol{X}_i\Omega_i^{-1}(\boldsymbol{\Theta})\boldsymbol{X}_i^{\intercal}\right)^{-1}\left\{\sum_{i=1}^{n} \boldsymbol{X}_i\Omega_i^{-1}(\boldsymbol{\Theta})\left(\frac{\partial^2\Omega_i(\boldsymbol{\Theta})}{\partial\theta\partial\theta'} - 2\frac{\partial\Omega_i(\boldsymbol{\Theta})}{\partial\theta}\Omega_i^{-1}(\boldsymbol{\Theta})\frac{\partial\Omega_i(\boldsymbol{\Theta})}{\partial\theta'}\right)\Omega_i(\boldsymbol{\Theta})^{-1}\boldsymbol{X}_i^{\intercal}\right.$$

$$\left.+ \left(\sum_{i=1}^{n} \boldsymbol{X}_i\Omega_i^{-1}(\boldsymbol{\Theta})\frac{\partial\Omega_i(\boldsymbol{\Theta})}{\partial\theta}\Omega_i(\boldsymbol{\Theta})^{-1}\boldsymbol{X}_i^{\intercal}\right)\left(\sum_{i=1}^{n} \boldsymbol{X}_i\Omega_i^{-1}(\boldsymbol{\Theta})\boldsymbol{X}_i^{\intercal}\right)\left(\sum_{i=1}^{n} \boldsymbol{X}_i\Omega_i^{-1}(\boldsymbol{\Theta})\frac{\partial\Omega_i(\boldsymbol{\Theta})}{\partial\theta'}\Omega_i(\boldsymbol{\Theta})^{-1}\boldsymbol{X}_i^{\intercal}\right)\right\}\right)$$

$$+ \sum_{i=1}^{n}\left(-\frac{1}{2}tr\left(\Omega_i(\boldsymbol{\Theta})^{-1}\frac{\partial\Omega_i(\boldsymbol{\Theta})}{\partial\theta'}\Omega_i(\boldsymbol{\Theta})^{-1}\frac{\partial\Omega_i(\boldsymbol{\Theta})}{\partial\theta}\right) - \frac{\partial\mu(\boldsymbol{\Theta}, \boldsymbol{X}_i)}{\partial\theta}\Omega_i(\boldsymbol{\Theta})^{-1}\frac{\partial\mu(\boldsymbol{\Theta}, \boldsymbol{X}_i)}{\partial\theta'}^{\intercal}\right)$$

This is what `information` output when the argument `type.information` is set to `"expected"` (up to a factor -1).

---

[3]if one is relative to the mean and the other to the variance then they are respectively θ and θ'

A.4 Degrees of freedom

Degrees of freedom are computed using a Satterthwaite approximation, i.e. for an estimate coefficient $\widehat{\beta} \in \widehat{\boldsymbol{\Theta}}$ with standard error $\sigma_{\widehat{beta}}$, the degree of freedom is:

$$df\left(\sigma_{\widehat{\beta}}\right) = \frac{2\sigma_{\widehat{\beta}}}{\mathbb{V}ar\left[\widehat{\sigma}_{\widehat{\beta}}\right]}$$

Using a first order Taylor expansion we can approximate the variance term as:

$$\mathbb{V}ar\left[\widehat{\sigma}_{\widehat{\beta}}\right] \approx \frac{\partial\widehat{\sigma}_{\widehat{\beta}}}{\partial\boldsymbol{\Theta}}\Sigma_{\boldsymbol{\Theta}}\frac{\partial\widehat{\sigma}_{\widehat{\beta}}}{\partial\boldsymbol{\Theta}}^{\intercal}$$

$$\approx c_{\beta}\left(\widehat{\mathcal{I}}_{\widehat{\boldsymbol{\Theta}}}\right)^{-1}\frac{\partial\widehat{\mathcal{I}}_{\widehat{\boldsymbol{\Theta}}}}{\partial\boldsymbol{\Theta}}\left(\widehat{\mathcal{I}}_{\widehat{\boldsymbol{\Theta}}}\right)^{-1}c_{\beta}^{\intercal}\Sigma_{\boldsymbol{\Theta}}c_{\beta}^{\intercal}\left(\widehat{\mathcal{I}}_{\widehat{\boldsymbol{\Theta}}}\right)^{-1}\frac{\partial\widehat{\mathcal{I}}_{\widehat{\boldsymbol{\Theta}}}}{\partial\boldsymbol{\Theta}}^{\intercal}\left(\widehat{\mathcal{I}}_{\widehat{\boldsymbol{\Theta}}}\right)^{-1}c_{\beta}$$

where $\Sigma_{\boldsymbol{\Theta}}$ is the variance-covariance matrix of all model coefficients, $\mathcal{I}_{\boldsymbol{\Theta}}$ the information matrix for all model coefficients, c_{β} a matrix used to select the element relative to β in the first derivative of the information matrix, and $\frac{\partial}{\partial\boldsymbol{\Theta}}$ denotes the vector of derivatives with respect to all model coefficients.

Appendix B Likelihood ratio test with the REML criterion

The blue term of Equation A in the log-likelihood is invariant to re-parameterisation while the red term is not. This means that a re-parametrisation of X into $\tilde{X} = BX$ with B invertible would not change the likelihood when using ML but would decrease the log-likelihood by $\log(|B|)$ when using REML.

Let's take an example:

```
## data(gastricbypassL, package = "LMMstar")
dfTest <- gastricbypassL
dfTest$glucagon2 <- dfTest$glucagon*2
```

where we multiply one column of the design matrix by 2. As mentionned previously this does not affect the log-likelihood when using ML:

```
logLik(lmm(weight ~ glucagon, data = dfTest, structure = UN(~time|id), method = "ML"))
logLik(lmm(weight ~ glucagon2, data = dfTest, structure = UN(~time|id), method = "ML"))
```

```
[1] -245.7909
[1] -245.7909
```

but it does when using REML:

```
logLik(lmm(weight ~ glucagon, data = dfTest, structure = UN(~time|id), method = "REML"))
logLik(lmm(weight ~ glucagon2, data = dfTest, structure = UN(~time|id), method = "REML"))
log(2)
```

```
[1] -245.0382
[1] -245.7313
[1] 0.6931472
```

Therefore, when comparing models with different mean effects there is a risk that the difference (or part of it) in log-likelihood is due to a new parametrisation and no only to a difference in model fit. This would typically be the case when adding an interaction where we can have a smaller restricted log-likehood when considering a more complex model:

```
set.seed(10)
dfTest$ff <- rbinom(NROW(dfTest), size = 1, prob = 0.5)
logLik(lmm(weight ~ glucagon, data = dfTest, structure = UN(~time|id), method = "REML"))
logLik(lmm(weight ~ glucagon*ff, data = dfTest, structure = UN(~time|id), method = "REML"))
```

```
[1] -245.0382
[1] -239.2056
```

This is quite counter-intuitive as more complex model should lead to better fit and would never happen when using ML:

```
logLik(lmm(weight ~ glucagon, data = dfTest, structure = UN(~time|id), method = "ML"))
logLik(lmm(weight ~ glucagon*ff, data = dfTest, structure = UN(~time|id), method = "ML"))
```

```
[1] -245.7909
[1] -237.3642
```

This is why, unless one knows what he/she is doing, it is not recommanded to use likelihood ratio test to assess relevance of mean parameters in mixed models estimated with REML.